

Gradient Boosting

David Rosenberg

New York University

March 23, 2016

Review: AdaBoost and FSAM

Adaptive Basis Function Model

- AdaBoost produces a classification score function of the form

$$\sum_{m=1}^M \alpha_m G_m(x)$$

- each G_m is a **weak classifier**
- The G_m 's are like basis functions, but they are learned from the data.
- Let's move beyond classification models...

Adaptive Basis Function Model

- **Base hypothesis space \mathcal{F}**
 - the “weak classifiers” in boosting context
- An **adaptive basis function expansion** over \mathcal{F} is

$$f(x) = \sum_{m=1}^M \alpha_m h_m(x),$$

- $h_m \in \mathcal{F}$ chosen in a learning process (“adaptive”)
 - $\alpha_m \in \mathbf{R}$ are **expansion coefficients**.
- **Note:** We are taking linear combination of outputs of $h_m(x)$.
 - Functions in $h_m \in \mathcal{F}$ must produce values in \mathbf{R} (or a vector space)

How to fit an adaptive basis function model?

- **Loss function:** $\ell(y, \hat{y})$
- **Base hypothesis space:** \mathcal{F} of **real-valued** functions
- Want to find

$$f(x) = \sum_{m=1}^M \alpha_m h_m(x)$$

that **minimizes empirical risk**

$$\frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- We'll proceed in stages, adding a new h_m in every stage.

Forward Stagewise Additive Modeling (FSAM)

- Start with $f_0 \equiv 0$.
- After $m-1$ stages, we have

$$f_{m-1} = \sum_{i=1}^{m-1} \nu_i h_i,$$

where $h_1, \dots, h_{m-1} \in \mathcal{F}$.

- Want to find
 - **step direction** $h_m \in \mathcal{F}$ and
 - **step size** $\nu_i > 0$
- So that

$$f_m = f_{m-1} + \nu_i h_m$$

minimizes empirical risk.

Forward Stagewise Additive Modeling

- 1 Initialize $f_0(x) = 0$.
- 2 For $m = 1$ to M :
 - 1 Compute:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathcal{F}} \sum_{i=1}^n \ell \left(y_i, f_{m-1}(x_i) + \underbrace{\nu h(x_i)}_{\text{new piece}} \right).$$

- 2 Set $f_m = f_{m-1} + \nu_m h$.
- 3 Return: f_M .

Example 1: Exponential Loss & Classifiers (AdaBoost)

- Loss function: $\ell(y, f(x)) = \exp(-yf(x))$.
- Base hypothesis space: $\mathcal{F} = \{h(x) : \mathcal{X} \rightarrow \{-1, 1\}\}$ (weak classifiers)
- Then Forward Stagewise Additive Modeling (FSAM) reduces to AdaBoost.
 - (See HTF Section 10.4 for proof.)

Example 2: Square Loss & Regression (L_2 -Boosting)

- Loss function: $\ell(y, f(x)) = (y - f(x))^2$
- Base hypothesis space: $\mathcal{F} = \{h(x) : \mathcal{X} \rightarrow \mathbf{R}\}$ (real-valued functions)
- Key step is

$$\begin{aligned}
 & \min_{v \in \mathbf{R}, h \in \mathcal{F}} \sum_{i=1}^n \left(y_i - \left[f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2 \\
 &= \min_{v \in \mathbf{R}, h \in \mathcal{F}} \sum_{i=1}^n \left(\underbrace{y_i - f_{m-1}(x_i)}_{\text{residual}} - vh(x_i) \right)^2
 \end{aligned}$$

Example 2: Square Loss & Regression (L_2 -Boosting)

- Simplifying assumption: \mathcal{F} is closed under scalar multiplication:
 - If $h \in \mathcal{F}$ then $ch \in \mathcal{F}$ for all $c \in \mathbf{R}$.
- Then step size is absorbed into \mathcal{F} we can just compute

$$\min_{h \in \mathcal{F}} \sum_{i=1}^n \left(\underbrace{y_i - f_{m-1}(x_i)}_{\text{residual}} - h(x_i) \right)^2$$

- This is square-loss regression on $(x_1, r_1), \dots, (x_n, r_n)$, where

$$r_i = y_i - f_{m-1}(x_i).$$

- **[Not linear regression** unless \mathcal{F} comprises linear functions.]
- This is called **L_2 -Boosting**.

FSAM: More Examples?

- The challenge with FSAM is solving

$$\min_{\nu \in \mathbf{R}, h \in \mathcal{F}} \sum_{i=1}^n \ell \left(y_i, f_{m-1}(x_i) + \underbrace{\nu h(x_i)}_{\text{new piece}} \right).$$

- Possibilities so far:
 - solve it exactly (e.g. AdaBoost)
 - reduce it to regression (e.g. L_2 -Boosting).
- But finding minimizer is not always easy for arbitrary
 - loss function and
 - base hypothesis space

Coordinate Descent Method

Coordinate Descent Method

Goal: Minimize $L(w) = L(w_1, \dots, w_d)$ over $w = (w_1, \dots, w_d) \in \mathbb{R}^d$.

- **Initialize** $w^{(0)} = 0$
- **while** not converged:
 - Choose a coordinate $j \in \{1, \dots, d\}$
 - $w_j^{\text{new}} \leftarrow \arg \min_{w_j} L(w_1^{(t)}, \dots, w_{j-1}^{(t)}, w_j, w_{j+1}^{(t)}, \dots, w_d^{(t)})$
 - $w^{(t+1)} \leftarrow w^{(t)}$
 - $w_j^{(t+1)} \leftarrow w_j^{\text{new}}$
 - $t \leftarrow t + 1$

FSAM as Coordinate Descent

- Suppose $\mathcal{F} = \{h_1, \dots, h_N\}$.
- Then $f_m = w_1 h_1 + \dots + w_N h_N$.
- Represent f_m by parameter vector $w^{(m)} = (w_1, \dots, w_N)$.
- Start with $w^{(0)} = 0$.
- After $m-1$ stages, we have $w^{(m-1)} = (w_1, \dots, w_N)$.
- Suppose m th step chooses
 - $h_m = h_j \in \mathcal{F}$ and $v_m \in \mathbf{R}$.
- Then $w^{(m)} = (w_1, \dots, w_{j-1}, w_j + v_m, w_{j+1}, \dots, w_N)$

Gradient Boosting / “Anyboost”

FSAM Looks Like Iterative Optimization

- The FSAM step

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbf{R}, h \in \mathcal{F}} \sum_{i=1}^n \ell \left(y_i, f_{m-1}(x_i) + \underbrace{\nu h(x_i)}_{\text{new piece}} \right).$$

- Hard part: finding the **best step direction** h .
- What if we looked for the **locally best** step direction?
 - like in gradient descent
- Approach:
 - Choose h_m to be something like a gradient in function space.
 - Roughly speaking, it will be the functional gradient projected onto \mathcal{F} .

Functional Gradient Descent: Main Idea

- We want to minimize

$$\sum_{i=1}^n \ell(y_i, f(x_i)).$$

- Take functional gradient w.r.t. f .
- Find function $h \in \mathcal{F}$ closest to gradient.
- Take a step in this "projected gradient" direction h .

"Functional" Gradient Descent

- We want to minimize

$$\sum_{i=1}^n \ell(y_i, f(x_i)).$$

- Only depends on f at the n training points.
- Define

$$\mathbf{f} = (f(x_1), \dots, f(x_n))^T$$

and write the objective function as

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, \mathbf{f}_i).$$

Functional Gradient Descent: Unconstrained Step Direction

- Consider gradient descent on

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, \mathbf{f}_i).$$

- The **negative gradient step direction** at \mathbf{f} is

$$-\mathbf{g} = -\nabla_{\mathbf{f}} J(\mathbf{f}),$$

which we can easily calculate.

Functional Gradient Descent: Projection Step

- Unconstrained step direction is

$$-\mathbf{g} = -\nabla_{\mathbf{f}} J(\mathbf{f}).$$

- Suppose \mathcal{F} is our weak hypothesis space.
- Find $h \in \mathcal{F}$ that is closest to $-\mathbf{g}$ at the training points, in the ℓ^2 sense:

$$\min_{h \in \mathcal{F}} \sum_{i=1}^n (-\mathbf{g}_i - h(x_i))^2.$$

- This is a least squares regression problem.
- \mathcal{F} should have **real-valued** functions.
- So the h that best approximates $-\mathbf{g}$ is our step direction.

Functional Gradient Descent: Step Size

- Finally, we choose a stepsize.
- Option 1 (Line search):

$$\nu_m = \arg \min_{\nu > 0} \sum_{i=1}^n \ell\{y_i, f_{m-1}(x_i) + \nu h_m(x_i)\}.$$

- Option 2: (Shrinkage parameter)
 - We consider $\nu = 1$ to be the full gradient step.
 - Choose a fixed $\nu \in (0, 1)$ – called a **shrinkage parameter**.
 - A value of $\nu = 0.1$ is typical – optimize as a hyperparameter .

The Gradient Boosting Machine

1 Initialize $f_0(x) = 0$.

2 For $m = 1$ to M :

1 Compute:

$$\mathbf{g}_m = \left(\left. \frac{\partial}{\partial f(x_i)} \left(\sum_{i=1}^n \ell\{y_i, f(x_i)\} \right) \right|_{f(x_i)=f_{m-1}(x_i)} \right)_{i=1}^n$$

2 Fit regression model to $-\mathbf{g}_m$:

$$h_m = \arg \min_{h \in \mathcal{F}} \sum_{i=1}^n ((-\mathbf{g}_m)_i - h(x_i))^2.$$

3 Choose fixed step size $\nu_m = \nu \in (0, 1]$, or take

$$\nu_m = \arg \min_{\nu > 0} \sum_{i=1}^n \ell\{y_i, f_{m-1}(x_i) + \nu h_m(x_i)\}.$$

4 Take the step:

$$f_m(x) = f_{m-1}(x) + \nu_m h_m(x)$$

The Gradient Boosting Machine: Recap

- Take any [sub]differentiable loss function.
- Choose a weak hypothesis space for regression.
- Choose number of steps (or a stopping criterion).
- Choose step size methodology.
- Then you're good to go!

Gradient Tree Boosting

Gradient Tree Boosting

- Common form of gradient boosting machine takes

$$\mathcal{F} = \{\text{regression trees of size } J\},$$

where J is the number of terminal nodes.

- $J = 2$ gives decision stumps
- HTF recommends $4 \leq J \leq 8$.
- Software packages:
 - Gradient tree boosting is implemented by the **gbm package** for R
 - as `GradientBoostingClassifier` and `GradientBoostingRegressor` in **sklearn**
- For trees, there are other tweaks on the algorithm one can do
 - See HTF 10.9-10.12 and