# Bayesian Regression

David Rosenberg

New York University

April 27, 2016

# Bayesian Statistics: Recap

## The Bayesian Method

1. Define the model:
   - Choose a probability model or "**likelihood model**":

     $$\{p(\mathcal{D} \mid \theta) \mid \theta \in \Theta\}.$$

   - Choose a distribution $p(\theta)$, called the **prior distribution**.

2. After observing $\mathcal{D}$, compute the **posterior distribution** $p(\theta \mid \mathcal{D})$.

3. Choose **action** based on $p(\theta \mid \mathcal{D})$.
   - e.g. $\mathbb{E}[\theta \mid \mathcal{D}]$ as point estimate for $\theta$
   - e.g. interval $[a, b]$, where $p(\theta \in [a, b] \mid \mathcal{D}) = 0.95$

## The Posterior Distribution

- By Bayes rule, can write the posterior distribution as

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta)p(\theta)}{p(\mathcal{D})}.$$

- **likelihood:** $p(\mathcal{D} \mid \theta)$
- **prior:** $p(\theta)$
- **marginal likelihood:** $p(\mathcal{D})$.
- Note: $p(\mathcal{D})$ is just a normalizing constant for $p(\theta \mid \mathcal{D})$. Can write

$$\underbrace{p(\theta \mid \mathcal{D})}_{\text{posterior}} \propto \underbrace{p(\mathcal{D} \mid \theta)}_{\text{likelihood}} \underbrace{p(\theta)}_{\text{prior}}.$$

## Summary

- Prior represents belief about θ before observing data $\mathcal{D}$.

- Posterior represents the **rationally "updated" beliefs** after seeing $\mathcal{D}$.

- All inferences and action-taking are based on posterior distribution.

# Bayesian Gaussian Linear Regression

# Bayesian Conditional Models

- Input space $\mathcal{X} = \mathbf{R}^d$      Output space $\mathcal{Y} = \mathbf{R}$

- **Conditional probability model**, or **likelihood model**:

$$\{p(y \mid x, \theta) \mid \theta \in \Theta\}$$

- Conditional here refers to the conditioning on the input $x$.
  - $x$'s are not governed by our probability model.
  - Everything conditioned on $x$ means "$x$ is known"

- **Prior distribution**: $p(\theta)$ on $\theta \in \Theta$

# Gaussian Regression Model

- Input space $\mathcal{X} = \mathbf{R}^d$    Output space $\mathcal{Y} = \mathbf{R}$
- **Conditional probability model**, or **likelihood model**:

$$y \mid x, w \quad \sim \quad \mathcal{N}\left(w^T x, \sigma^2\right),$$

  for some known $\sigma^2 > 0$.
- **Parameter space?** $\mathbf{R}^d$.
- **Data:** $\mathcal{D} = \{(x_1, y_1), \ldots, (x_n, y_n)\}$
    - **Notation**: $y = (y_1, \ldots, y_n)$ and $x = (x_1, \ldots, x_n)$.
    - Assume $y_i$'s are **conditionally independent**, given $x$ and $w$.

# Gaussian Likelihood

- The **likelihood** of $w \in \mathbf{R}^d$ for the data $\mathcal{D}$ is

$$
\begin{aligned}
p(y \mid x, w) &= \prod_{i=1}^{n} p(y_i \mid x_i, w) \qquad \text{by conditional independence.} \\
&= \prod_{i=1}^{n} \left[ \frac{1}{\sigma \sqrt{2\pi}} \exp\left( -\frac{(y_i - w^T x_i)^2}{2\sigma^2} \right) \right]
\end{aligned}
$$

- You should see **in your head**[1] that the **MLE** is

$$
\begin{aligned}
w_{\text{MLE}}^* &= \underset{w \in \mathbf{R}^d}{\arg\max}\, p(y \mid x, w) \\
&= \underset{w \in \mathbf{R}^d}{\arg\min} \sum_{i=1}^{n} (y_i - w^T x_i)^2.
\end{aligned}
$$

---

[1]See https://davidrosenberg.github.io/ml2015/docs/8.Lab.glm.pdf, slide 5.

## Priors and Posteriors

- Choose a Gaussian **prior distribution** $p(w)$ on $\mathbf{R}^d$:

$$w \sim \mathcal{N}(0, \Sigma_0)$$

for some **covariance matrix** $\Sigma_0 \succ 0$ (i.e. $\Sigma_0$ is spd).

- **Posterior distribution**

$$
\begin{aligned}
p(w \mid \mathcal{D}) &= p(w \mid x, y) \\
&= p(y \mid x, w) \, p(w) / p(y) \\
&\propto p(y \mid x, w) p(w) \\
&= \prod_{i=1}^{n} \left[ \frac{1}{\sigma \sqrt{2\pi}} \exp\left( -\frac{(y_i - w^T x_i)^2}{2\sigma^2} \right) \right] \text{ (likelihood)} \\
&\quad \times |2\pi\Sigma_0|^{-1/2} \exp\left( -\frac{1}{2} w^T \Sigma_0^{-1} w \right) \text{ (prior)}
\end{aligned}
$$

# What does the posterior give us?

- **Likelihood model**: $y \mid x, w \sim \mathcal{N}\left(w^T x, \sigma^2\right)$
- **Prior distribution:** $w \sim \mathcal{N}(0, \Sigma_0)$
- Given data, compute **posterior distribution**: $p(w \mid \mathcal{D})$.
- If we knew $w$, best prediction function (for square loss) is

$$\hat{y}(x) = \mathbb{E}[y \mid x, w] = w^T x.$$

- Prior $p(w)$ and posterior $p(w \mid \mathcal{D})$

  - give **distributions over prediction functions**!

# Gaussian Regression Example

# Example in 1-Dimension: Setup

- Input space $\mathcal{X} = [-1, 1]$    Output space $\mathcal{Y} = \mathbf{R}$
- Let's suppose for any $x$, $y$ is generated as

$$y = w_0 + w_1 x + \varepsilon,$$

where $\varepsilon \sim \mathcal{N}(0, 0.2^2)$.

- Written another way, the **likelihood model** is

$$y \mid x, w_0, w_1 \quad \sim \quad \mathcal{N}\left(w_0 + w_1 x, \, 0.2^2\right).$$

- What's the parameter space? $\mathbf{R}^2$.
- **Prior distribution:** $w = (w_0, w_1) \sim \mathcal{N}\left(0, \frac{1}{2} I\right)$
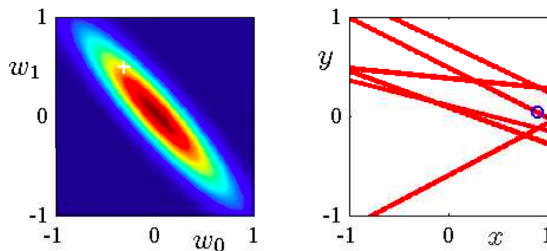
# Example in 1-Dimension: Prior Situation

- **Prior distribution:** $w = (w_0, w_1) \sim \mathcal{N}\left(0, \frac{1}{2}I\right)$
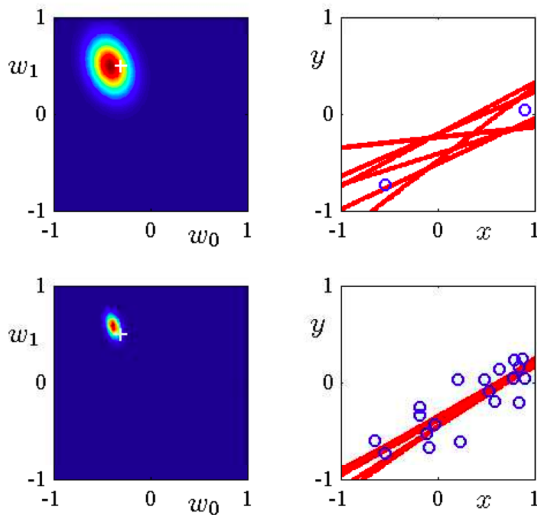


- On right, $y(x) = \mathbb{E}[y \mid x, w]$, for randomly chosen $w \sim \mathcal{N}\left(0, \frac{1}{2}I\right)$.
  - $y(x) = w_0 + w_1 x$ for random $(w_0, w_1) \sim p(\theta) = \mathcal{N}(0, \frac{1}{2}I)$.

Bishop's PRML Fig 3.7

# Example in 1-Dimension: 1 Observation



- On left: posterior distribution; white '+' indicates true parameters
- On right: blue circle indicates the training observation

# Example in 1-Dimension: 2 and 20 Observations



Bishop's PRML Fig 3.7

# Gaussian Regression Continued

# Closed Form for Posterior

- Model:

$$
\begin{aligned}
w &\sim \quad \mathcal{N}(0, \Sigma_0) \\
y_i \,|\, x, w &\quad \text{i.i.d.} \quad \mathcal{N}(w^T x_i, \sigma^2)
\end{aligned}
$$

- Design matrix $X$;        Response column vector $y$
- **Posterior distribution is a Gaussian distribution:**

$$
\begin{aligned}
w \,|\, \mathcal{D} &\sim \quad \mathcal{N}(\mu_P, \Sigma_P) \\
\mu_{\mathbf{P}} &= \quad \left(X^T X + \sigma^2 \Sigma_0^{-1}\right)^{-1} X^T y \\
\Sigma_{\mathbf{P}} &= \quad \left(\sigma^{-2} X^T X + \Sigma_0^{-1}\right)^{-1}
\end{aligned}
$$

- **Posterior Variance $\Sigma_P$ gives us a natural uncertainty measure.**

See Rasmussen and Williams' *Gaussian Processes for Machine Learning*, Ch 2.1.
http://www.gaussianprocess.org/gpml/chapters/RW2.pdf

## Closed Form for Posterior

- **Posterior distribution is a Gaussian distribution:**

$$
\begin{aligned}
w \mid \mathcal{D} &\sim \mathcal{N}(\mu_P, \Sigma_P) \\
\mu_P &= \left(X^T X + \sigma^2 \Sigma_0^{-1}\right)^{-1} X^T y \\
\Sigma_P &= \left(\sigma^{-2} X^T X + \Sigma_0^{-1}\right)^{-1}
\end{aligned}
$$

- The **MAP estimator** and the **posterior mean** are given by

$$
\mu_P = \left(X^T X + \sigma^2 \Sigma_0^{-1}\right)^{-1} X^T y
$$

- For the prior variance $\Sigma_0 = \frac{\sigma^2}{\lambda} I$, we get

$$
\mu_P = \left(X^T X + \lambda I\right)^{-1} X^T y,
$$

which is of course the ridge regression solution.

# Posterior Mean and Posterior Mode (MAP)

- Posterior density for $\Sigma_0 = \frac{\sigma^2}{\lambda} I$:

$$p(w \mid \mathcal{D}) \quad \propto \quad \underbrace{\exp\left(-\frac{\lambda}{2\sigma^2}\|w\|^2\right)}_{\text{prior}} \underbrace{\prod_{i=1}^{n} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)}_{\text{likelihood}}$$

- To find MAP, sufficient to minimize the negative log posterior:

$$\begin{aligned}
\hat{w}_{\text{MAP}} &= \underset{w \in \mathbf{R}^d}{\arg\min}\left[-\log p(w \mid \mathcal{D})\right] \\
&= \underset{w \in \mathbf{R}^d}{\arg\min} \underbrace{\sum_{i=1}^{n}(y_i - w^T x_i)^2}_{\text{log-likelihood}} + \underbrace{\lambda\|w\|^2}_{\text{log-prior}}
\end{aligned}$$

- Which is the ridge regression objective.

# Predictive Distribution

- Given a new input point $x_{\textbf{new}}$, how to predict $y_{\textbf{new}}$ ?
- **Predictive distribution**

$$
\begin{aligned}
p(y_{\textbf{new}} \mid x_{\textbf{new}}, \mathcal{D}) &= \int p(y_{\textbf{new}} \mid x_{\textbf{new}}, \theta, \mathcal{D}) p(\theta \mid \mathcal{D}) \, d\theta \\
&= \int p(y_{\textbf{new}} \mid x_{\textbf{new}}, \theta) p(\theta \mid \mathcal{D}) \, d\theta
\end{aligned}
$$

- For Gaussian regression, predictive distribution has closed form.

# Closed Form for Predictive Distribution

- **Model**:

$$
\begin{aligned}
w &\sim \mathcal{N}(0, \Sigma_0) \\
y_i \mid x, w \quad \text{i.i.d.} &\quad \mathcal{N}(w^T x_i, \sigma^2)
\end{aligned}
$$

- **Predictive Distribution**

$$
p(y_{\text{new}} \mid x_{\text{new}}, \mathcal{D}) = \int p(y_{\text{new}} \mid x_{\text{new}}, w) p(w \mid \mathcal{D}) \, d\theta.
$$

  - Averages over prediction for each $\theta$, weighted by posterior distribution.

- **Closed form**:

$$
\begin{aligned}
y_{\text{new}} \mid x_{\text{new}}, \mathcal{D} &\sim \mathcal{N}(\eta_{\text{new}}, \sigma_{\text{new}}) \\
\eta_{\text{new}} &= \mu_P^T x_{\text{new}} \\
\sigma_{\text{new}} &= \underbrace{x_{\text{new}}^T \Sigma_P x_{\text{new}}}_{\text{from variance in } \theta} + \underbrace{\sigma^2}_{\text{inherent variance in } y}
\end{aligned}
$$

# Predictive Distributions

- With predictive distributions, can give mean prediction with error bands:



Rasmussen and Williams' *Gaussian Processes for Machine Learning*, Fig.2.1(b)